✦ HACKATHON 2026 OFFICIAL DECK

INKLOOM AI

Intelligent Brand Governance Platform

"Create confidently. Stay consistent. Protect your brand."

AUTOMATED GOVERNANCE

Multi-layer brand compliance powered by Google Gemini 2.5 Flash.

ZERO HALLUCINATION

Strict factual repository check flags unverified claims.

HUMAN-IN-THE-LOOP

Certified compliance gate ensuring human editorial approval.

The Problem: Why Modern Brand Communication Breaks Down

Organizations publish content through dozens of employees, agencies, copywriters, and autonomous generative AI tools. Without centralized governance, teams accidentally release brand violations that erode public trust and invite regulatory risk.

1. Unapproved Claims & Fact Drift

HIGH LEGAL RISK

Teams accidentally publish outdated statistics, unverified cure rates, or speculative financial warranties that trigger regulatory fines and public

2. Inconsistent Tone of Voice

BRAND IDENTITY DECAY

One department sounds sober and institutional while an external agency posts meme-heavy colloquialisms, confusing enterprise buyers and degrading brand prestige.

3. Prohibited & Off-Brand Terminology

COMPLIANCE VIOLATION

Misuse of forbidden nomenclature (e.g. 'AI Doctor' instead of 'Clinical Decision Support'), violating trademark boundaries and medical/financial positioning.

4. Manual Editorial Bottlenecks

SEVERE INEFFICIENCY

Compliance officers drown in endless Slack pings and email threads, leading to either delayed campaign velocity or unchecked rogue publishing.

The Solution: Automated Governance with Explainable Oversight

Inkloom AI connects directly to verified organizational guidelines, approved facts, and prohibited claim libraries. Content drafts are automatically evaluated across 5 core dimensions before reaching human reviewers.

Tone & Voice Analysis

Calibrates formality, empathy, and reading grade level for target personas.

Brand Terminology

Identifies restricted words and recommends approved preferred nomenclature.

Factual Verification

Checks claims against verified facts; flags unsupported data as 'Requires

Messaging & Claims

Detects misleading warranties, speculative yield promises, or

Visual Brand Guidance

Cross-references brand imagery and logos against official visual

■ THE INKLOOM CONSISTENCY SCORE (0–100)

- Instant algorithmic index derived from weighted severity deductions (Critical: -20, Warning: -10, Suggestion: -5).
- Radial color-coded visual indicator with clear breakdown: Critical Violations, Warnings, Suggestions & Passed Checks.
- Enterprise benchmark: Submissions below 70% threshold are gated from publication until corrections are applied.

Key Features: Explainable Diagnostics & One-Click Suggestion Engine

Inkloom AI bridges the gap between raw AI generation and audit-ready corporate compliance.

EXPLAINABLE AI DIAGNOSTICS

Every flagged issue includes:

- Exact flagged text snippet highlighted
- Clear explanation of the violation
- Cited brand guideline & source document
- Severity classification (Critical / Warning)

ONE-CLICK SUGGESTION ENGINE

Contextual rewrite capabilities:

- Side-by-side comparison (Original vs Improved)
- Preserves marketing intent & brand voice
- 4 action buttons: Apply, Dismiss, Edit, Regenerate
- Replaces text seamlessly in draft

MULTI-FORMAT GUIDELINE REPOSITORY

Enterprise document ingestion:

- Upload PDF, DOCX, and TXT style guides
- Automatic text parsing & index extraction
- Active brand switching (Healthcare & FinTech)
- Verified facts & prohibited claim catalogs

Human-in-the-Loop Workflow: Certified Editorial Accountability

AI provides diagnostic guidance; human compliance officers make certified publishing decisions.

STEP 1

Draft Submitted

Creator pastes text or uploads draft across channels (email, social,

STEP 2

Gemini AI Audit

Scans text against brand voice, prohibited words, and approved facts in

STEP 3

Reviewer Queue

Flagged content routes to authorized Compliance Officer dashboard.

STEP 4

Human Decision

Officer inspects original vs improved text, adds notes, and issues

Official Human Determination States:

APPROVED

Content satisfies all guidelines and is certified for public distribution.

CHANGES REQUESTED

Specific violations flagged for creator revision before reconsideration.

REJECTED

Content fundamentally violates legal boundaries or brand positioning.

Technology Stack & Production Cloud Deployment

Engineered with modern, scalable, and resilient enterprise cloud infrastructure.

FRONTEND

Framework
React 18 + Vite

Styling
Tailwind CSS

Visuals
Recharts & Lucide Icons

State
React Context (Auth & Brand)

Routing
React Router DOM v6 SPA

Hosting
Render Static Site (Live)

BACKEND API

Runtime
Node.js (v24 LTS)

Framework
Express.js REST APIs

Database
MongoDB & Mongoose

Authentication
Firebase Auth + JWT

Document Engine
Multer Ingestion Gateway

Hosting
Render Web Service

AI SERVICE

Language
Python 3.14

Framework
FastAPI + Uvicorn ASGI

LLM Engine
Google Gemini 2.5 Flash

Doc Extract
PyPDF & python-docx

Fail-Safe
Deterministic Rule Engine

Output Mode
Strict JSON Schema

Business Value, ROI & Strategic Roadmap

78%

Faster Review Cycles

Eliminates manual editorial ping-pong with insta...

0%

Unsubstantiated Claims

Factual repository enforcement prevents deceptiv...

100%

Human Accountability

Every live release has certified, timestamped hu...

Strategic Product Roadmap:

- **Phase 1: Foundation (Current Hackathon MVP)**
  Full-stack SaaS, multi-brand repository, Gemini 2.5 Flash analysis, and human approval queue.
- **Phase 2: Omnichannel CMS Connectors**
  Direct publishing integrations for WordPress, HubSpot, Webflow, and LinkedIn API.
- **Phase 3: Real-Time Browser Extension Copilot**
  Chrome DevTools & sidecar extension flagging brand compliance directly inside Google Docs and Notion.
- **Phase 4: Multimodal Visual & Audio Governance**
  Detect logo safety margins, brand color deviations, and audio podcast voiceover compliance.